

From Plausibility Machines to Narrative Minds

Vision, Attention, and the Architecture of Artificial Cognition

Your Name

Your Institution

July 4, 2026

Outline

- 1 An Analogy: Human and Computer Vision
- 2 Chiang, Tufekci, and the Phenomenological Critique
- 3 “Attention Is All You Need”
- 4 Coglinks and Uncertainty
- 5 Time, Prediction, and Memory
- 6 Phenomenological Experience and Emotion
- 7 Narrative in Cognition
- 8 What Architecture Is Closer?
- 9 What Training Is Required?
- 10 How Might a Narrative-Processing AI Differ?
- 11 Notes / Parking Lot
- 12 Conclusion

An Analogy: Human and Computer Vision

- **Edge detection** — low-level feature extraction

An Analogy: Human and Computer Vision

- **Edge detection** — low-level feature extraction
- **Abstraction** — features aggregate into intermediate representations

Why this analogy

Raw input → abstraction → recognition is a useful template for thinking about cognitive architecture generally. Where does narrative fit in this pipeline?

An Analogy: Human and Computer Vision

- **Edge detection** — low-level feature extraction
- **Abstraction** — features aggregate into intermediate representations
- **Object recognition** — abstraction resolves into a labeled category

Why this analogy

Raw input → abstraction → recognition is a useful template for thinking about cognitive architecture generally. Where does narrative fit in this pipeline?

- Chiang and Tufekci: LLMs as **plausibility machines** (Chiang 2024; Tufekci 2026) — producing text that reads as coherent without underlying understanding
- The move from *simulating* consciousness to *having* consciousness is often treated as small, or even inevitable

- Chiang and Tufekci: LLMs as **plausibility machines** (Chiang 2024; Tufekci 2026) — producing text that reads as coherent without underlying understanding
- The move from *simulating* consciousness to *having* consciousness is often treated as small, or even inevitable

Claim

That leap is in fact massive, and unlikely.

The Phenomenological Critique

Central claim

There is no “what it is like” for a current AI system.

The Phenomenological Critique

Central claim

There is no “what it is like” for a current AI system.

- **What it is:** procedural circuits — comparable to functions in a CS sense

The Phenomenological Critique

Central claim

There is no “what it is like” for a current AI system.

- **What it is:** procedural circuits — comparable to functions in a CS sense
- **What is (partly) missing:** valence and salience

“Attention Is All You Need” (Vaswani et al. 2017)

- What attention consists of, for an LLM
- Stacked transformers and attention
- Loss minimization and static associative valuation

- Circuits imitated
- Predefined decision choices

Time, Prediction, and Memory

- Predictive processing vs. uncertainty
- Episodic and semantic memory

Open question

What would have to be true of a system for phenomenological experience and emotion — not merely their outward signatures — to be present?

- As active inference

- As **active inference**
- As an **abstraction scaffold**

What Architecture Is Closer?

- Self-salience models
- Less terrifying possibilities that simulate emotion

What Training Is Required?

- Direction 1: [To be developed]
- Direction 2: [To be developed]

How Might a Narrative-Processing AI Differ?

- 1 Not procedural?

How Might a Narrative-Processing AI Differ?

- ① Not procedural?
- ② Predictive processing (with respect to dynamic inputs)?

Thus

A closer semblance of self.

How Might a Narrative-Processing AI Differ?

- ① Not procedural?
- ② Predictive processing (with respect to dynamic inputs)?
- ③ Nested experiential hierarchies — including reflection on what experiences reveal about self-function?

Thus

A closer semblance of self.

- **Oh no Waymo —**

possible example of imitation-without-understanding in the wild

Waymos and other AIs harnessed to bodies perform processing operations related to moving through the world with particular sets of prescribed motivations. There's lots to question about the notion of 'motivations', which I hope to do in another forum. But focusing for now just on the problem of regulating a body moving in time through space, detecting environmental conditions is an ongoing, dynamic problem. At its core, it requires recognizing other agents and dynamic forces in the world. That is, slow feature analysis is a baseline for experiential processing (Song and Zhao 2022). It builds temporal parsing into experiencing. Slow feature analysis combines with valuative processing (salience and valence) that is itself dynamic and interactive. In terms of attention, a stabilization in valuation can be a catalyst of shifting attention, which marks temporal experience with subjective meaning.

Open questions: Can an experiencing AI have a semblance of these processes? How might it relate to the question of whether the AI has subjective experience?

Thank you!

Questions?

-  Chiang, Ted (2024). *PLACEHOLDER – Ted Chiang on LLMs as “plausibility machines”* (verify exact source). Placeholder citation — fill in exact title, publication, and date.
-  Miller, Brook (2024). *Narrativity in cognition*. Springer.
-  Song, Pengyu and Chunhui Zhao (2022). “Slow down to go better: A survey on slow feature analysis”. In: *IEEE Transactions on Neural Networks and Learning Systems* 35.3, pp. 3416–3436.
-  Tufekci, Zeynep (2026). “The Unstoppable Force of A.I. Hype Is Meeting One Immovable Fact”. In: *The New York Times*. Opinion, Guest Essay; published June 30, 2026.
-  Vaswani, Ashish et al. (2017). “Attention is all you need”. In: *Advances in Neural Information Processing Systems* 30.